

VisionThing

February 2, 2026

```
[1]: #psychometrics.ai
```

```
import os
from PIL import Image
```

```
[2]: #collect pngs and check they open
```

```
import glob

# Collect only PNGs from the current folder
image_paths = sorted(glob.glob("./.*.png"))[:30]

print(f"Using {len(image_paths)} PNGs:")
for p in image_paths:
    print(" -", p)

# Quick open-check to catch any corrupted files early
bad = []
for p in image_paths:
    try:
        with Image.open(p) as img:
            img.verify() # verifies file integrity
    except Exception as e:
        bad.append((p, str(e)))

if bad:
    print(" Some images failed integrity check:")
    for p, err in bad:
        print(f" - {p}: {err}")
    raise SystemExit("Please fix or remove the bad files before proceeding.")
```

Using 30 PNGs:

- ./A-1.png
- ./A-10.png
- ./A-11.png
- ./A-12.png
- ./A-13.png
- ./A-14.png

```
- ./A-15.png
- ./A-16.png
- ./A-17.png
- ./A-18.png
- ./A-19.png
- ./A-2.png
- ./A-20.png
- ./A-21.png
- ./A-22.png
- ./A-23.png
- ./A-24.png
- ./A-25.png
- ./A-26.png
- ./A-27.png
- ./A-28.png
- ./A-29.png
- ./A-3.png
- ./A-30.png
- ./A-4.png
- ./A-5.png
- ./A-6.png
- ./A-7.png
- ./A-8.png
- ./A-9.png
```

```
[3]: from sentence_transformers import SentenceTransformer
import torch

device = "cuda" if torch.cuda.is_available() else "cpu"

# IMPORTANT: use the correct model ID
model = SentenceTransformer("sentence-transformers/clip-ViT-B-32",  
    device=device)

print("Loaded:", type(model))
print("Device:", device)
```

```
/opt/anaconda3/lib/python3.11/site-
packages/sentence_transformers/cross_encoder/CrossEncoder.py:11:
TqdmExperimentalWarning: Using `tqdm.autonotebook.tqdm` in notebook mode. Use
`tqdm.tqdm` instead to force console mode (e.g. in jupyter console)
    from tqdm.autonotebook import tqdm, trange

Loaded: <class 'sentence_transformers.SentenceTransformer.SentenceTransformer'>
Device: cpu
```

```
[4]: import numpy as np
```

```

# Create embeddings for the 30 images
embeddings = model.encode(
    image_paths,                      # the list of your PNGs
    batch_size=8,                     # fine for small sets
    convert_to_numpy=True,
    normalize_embeddings=True, # good for cosine similarity
    show_progress_bar=True
)

print("Embeddings created.")
print("Shape:", embeddings.shape)  # should be (30, 512)

```

Batches: 0% | 0/4 [00:00<?, ?it/s]

Embeddings created.

Shape: (30, 512)

```

[5]: import pandas as pd

# image_paths = [...]      # already created earlier
# embeddings = [...]        # already computed earlier

df = pd.DataFrame(embeddings)
df.insert(0, "filename", image_paths)

# Show a scrollable full-width table inside Jupyter
from IPython.display import display, HTML
display(HTML(df.to_html(max_rows=200, max_cols=200, notebook=True)))

```

<IPython.core.display.HTML object>

```

[6]: from sklearn.metrics.pairwise import cosine_similarity

# Compute cosine similarity matrix
similarity_matrix = cosine_similarity(embeddings)

# Create labeled DataFrame
filenames = [os.path.basename(p) for p in image_paths]
sim_df = pd.DataFrame(similarity_matrix, columns=filenames)
sim_df.insert(0, "filename", filenames)

# Display
display(HTML(sim_df.to_html(max_rows=200, max_cols=200, notebook=True)))

```

<IPython.core.display.HTML object>

```

[7]: from sklearn.decomposition import FactorAnalysis
import matplotlib.pyplot as plt

```

```

# Extract just the similarity values (exclude filename column)
sim_matrix = sim_df.iloc[:, 1: ].values

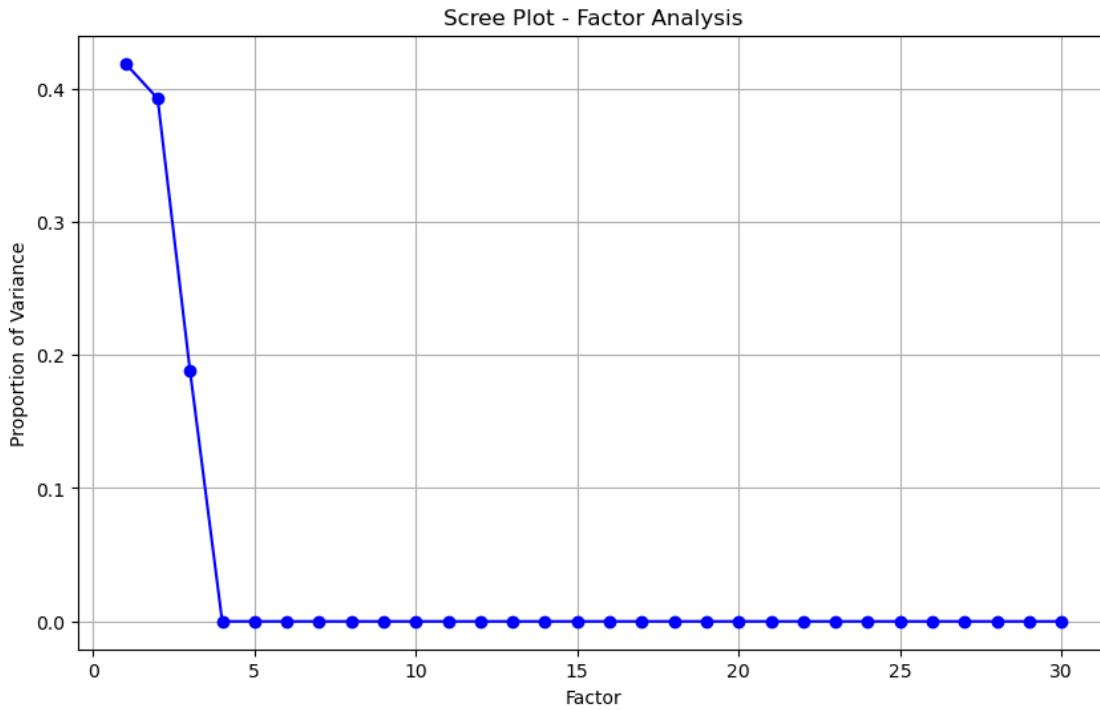
# Factor Analysis with Maximum Likelihood
fa = FactorAnalysis(n_components=None, rotation=None)
fa.fit(sim_matrix)

# Scree plot (using noise variance as proxy for explained variance)
# Note: FA doesn't have explained_variance_ratio_ like PCA
# We can plot the components' variances
component_vars = np.var(fa.transform(sim_matrix), axis=0)
total_var = np.sum(component_vars)
explained_var_ratio = component_vars / total_var

plt.figure(figsize=(10, 6))
plt.plot(range(1, len(explained_var_ratio) + 1),
         explained_var_ratio, 'bo-')
plt.xlabel('Factor')
plt.ylabel('Proportion of Variance')
plt.title('Scree Plot - Factor Analysis')
plt.grid(True)
plt.show()

# Get loadings with labels
loadings = fa.components_.T
loadings_df = pd.DataFrame(
    loadings,
    index=filenames,
    columns=[f'Factor{i+1}' for i in range(loadings.shape[1])]
)
display(HTML(loadings_df.to_html(max_rows=200, max_cols=200, notebook=True)))

```



<IPython.core.display.HTML object>

```
[8]: import pandas as pd

# Read the CSV file
form1 = pd.read_csv('form1.csv')

# Check what the columns are actually named
print("Column names:", form1.columns.tolist())
print("First few rows:")
print(form1.head())

# Assign using .values to avoid index mismatch
loadings_df['dis'] = form1.iloc[:, 0].values # First column (dis)
loadings_df['dif'] = form1.iloc[:, 1].values # Second column (dif)

# Display the updated DataFrame
display(HTML(loadings_df.to_html(max_rows=200, max_cols=200, notebook=True)))
```

Column names: ['dis', 'dif']

First few rows:

	dis	dif
0	1.11	1.15
1	1.10	0.94
2	1.28	-1.22

```
3 1.09 1.33
4 0.68 0.95
```

```
<IPython.core.display.HTML object>
```

```
[9]: # Correlations between factor 1 and dis and dif
```

```
corr_dif_factor1 = loadings_df['dif'].corr(loadings_df['Factor1'])
corr_dis_factor1 = loadings_df['dis'].corr(loadings_df['Factor1'])
print(f"dif vs Factor1: r = {corr_dif_factor1:.3f}")
print(f"dis vs Factor1: r = {corr_dis_factor1:.3f}")

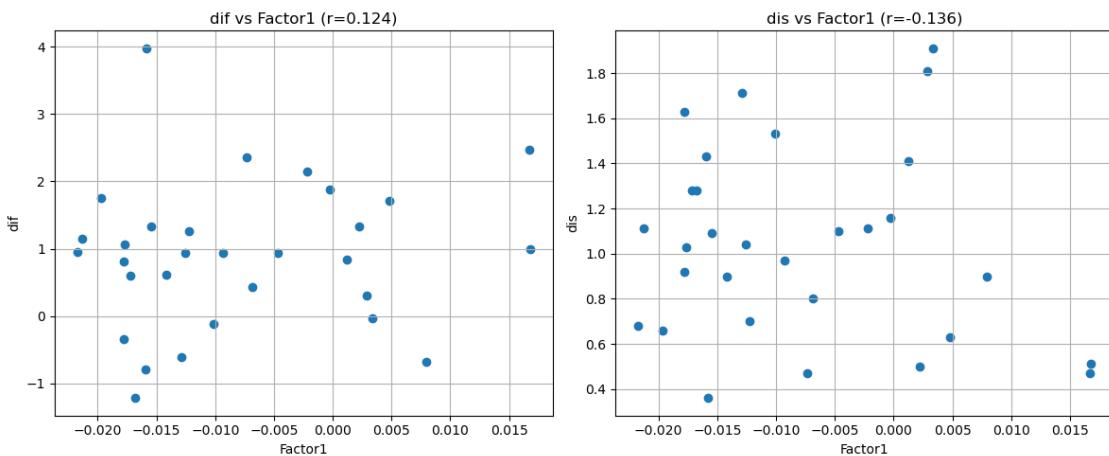
# Scatter plots
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

axes[0].scatter(loadings_df['Factor1'], loadings_df['dif'])
axes[0].set_xlabel('Factor1')
axes[0].set_ylabel('dif')
axes[0].set_title(f'dif vs Factor1 (r={corr_dif_factor1:.3f})')
axes[0].grid(True)

axes[1].scatter(loadings_df['Factor1'], loadings_df['dis'])
axes[1].set_xlabel('Factor1')
axes[1].set_ylabel('dis')
axes[1].set_title(f'dis vs Factor1 (r={corr_dis_factor1:.3f})')
axes[1].grid(True)

plt.tight_layout()
plt.show()
```

```
dif vs Factor1: r = 0.124
dis vs Factor1: r = -0.136
```



```
[10]: # Correlations between factor 2 and dis and dif

corr_dif_factor2 = loadings_df['dif'].corr(loadings_df['Factor2'])
corr_dis_factor2 = loadings_df['dis'].corr(loadings_df['Factor2'])
print(f"dif vs Factor2: r = {corr_dif_factor2:.3f}")
print(f"dis vs Factor2: r = {corr_dis_factor2:.3f}")

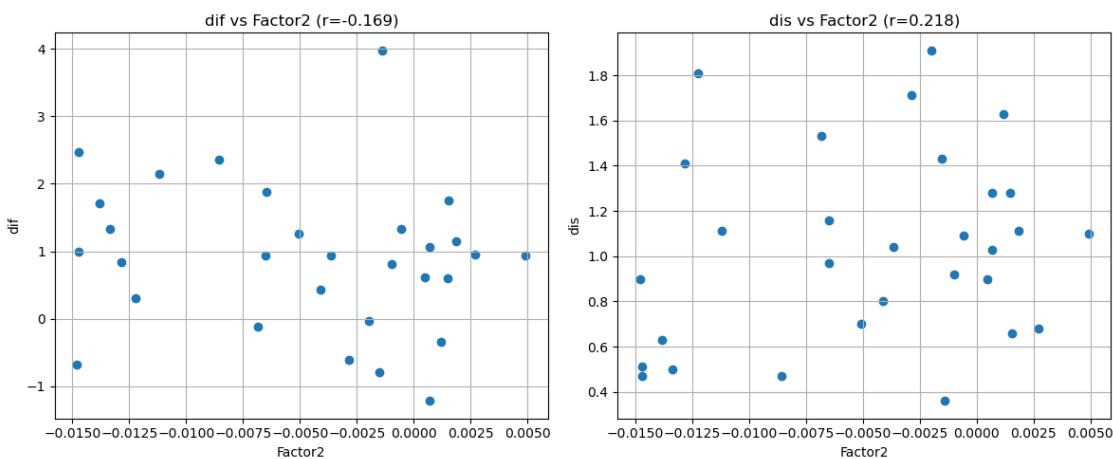
# Scatter plots
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

axes[0].scatter(loadings_df['Factor2'], loadings_df['dif'])
axes[0].set_xlabel('Factor2')
axes[0].set_ylabel('dif')
axes[0].set_title(f'dif vs Factor2 (r={corr_dif_factor2:.3f})')
axes[0].grid(True)

axes[1].scatter(loadings_df['Factor2'], loadings_df['dis'])
axes[1].set_xlabel('Factor2')
axes[1].set_ylabel('dis')
axes[1].set_title(f'dis vs Factor2 (r={corr_dis_factor2:.3f})')
axes[1].grid(True)

plt.tight_layout()
plt.show()
```

dif vs Factor2: r = -0.169
dis vs Factor2: r = 0.218



```
[11]: # Correlations between factor 3 and dis and dif
```

```
corr_dif_factor3 = loadings_df['dif'].corr(loadings_df['Factor3'])
```

```

corr_dis_factor3 = loadings_df['dis'].corr(loadings_df['Factor3'])
print(f"dif vs Factor3: r = {corr_dif_factor3:.3f}")
print(f"dis vs Factor3: r = {corr_dis_factor3:.3f}")

# Scatter plots
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

axes[0].scatter(loadings_df['Factor3'], loadings_df['dif'])
axes[0].set_xlabel('Factor3')
axes[0].set_ylabel('dif')
axes[0].set_title(f'dif vs Factor3 (r={corr_dif_factor3:.3f})')
axes[0].grid(True)

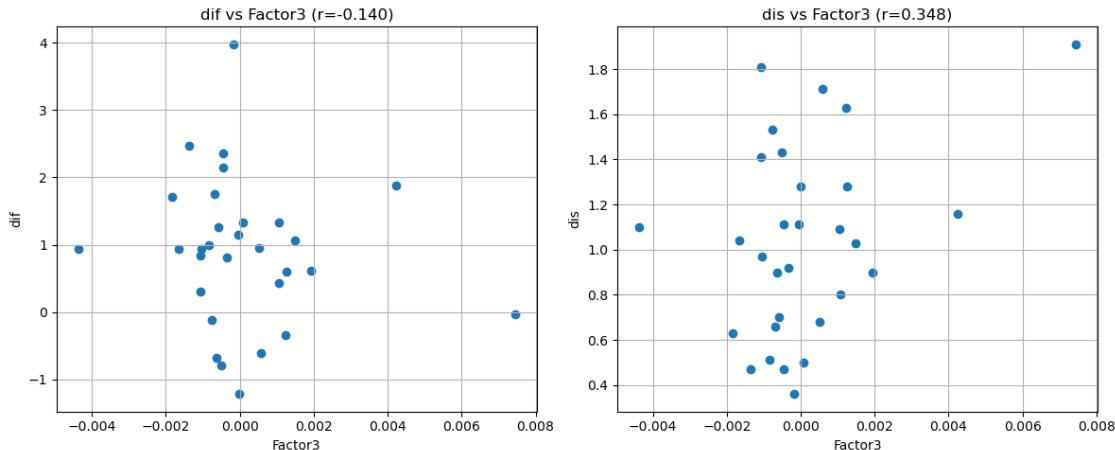
axes[1].scatter(loadings_df['Factor3'], loadings_df['dis'])
axes[1].set_xlabel('Factor3')
axes[1].set_ylabel('dis')
axes[1].set_title(f'dis vs Factor3 (r={corr_dis_factor3:.3f})')
axes[1].grid(True)

plt.tight_layout()
plt.show()

```

dif vs Factor3: r = -0.140

dis vs Factor3: r = 0.348



[12]: # --- Try PCA on raw embeddings + correlations with IRT params (dif, dis)

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler

```

```

from sklearn.decomposition import PCA
from scipy.stats import pearsonr

# ---- Gather inputs ----
# embeddings: np.ndarray shape (n_items, 512)
# filenames: list[str] length n_items
# dif/dis: arrays-like length n_items (fallback: read from loadings_df if
# ↵present)
def _arr(x):
    return np.asarray(x, dtype=float).reshape(-1)

n_items = embeddings.shape[0]
if 'filenames' not in globals():
    # Try to reconstruct from earlier sim_df or file paths if needed
    try:
        filenames = list(loadings_df.index)
    except Exception as _:
        filenames = [f"item_{i+1}" for i in range(n_items)]

# Pull dif/dis from globals or loadings_df if available
if 'dif' in globals() and 'dis' in globals():
    dif_vec, dis_vec = _arr(dif), _arr(dis)
else:
    try:
        dif_vec = _arr(loadings_df['dif'].values)
        dis_vec = _arr(loadings_df['dis'].values)
    except Exception as e:
        raise RuntimeError("Need vectors 'dif' and 'dis' (or loadings_df['dif']/
        ↵'dis'])") from e

assert len(dif_vec) == n_items and len(dis_vec) == n_items, "Length mismatch
    ↵for dif/dis vs embeddings."

# ---- PCA on raw embeddings ----
X = embeddings # (n_items, 512)
Xz = StandardScaler(with_mean=True, with_std=True).fit_transform(X)
pca = PCA(n_components=min(10, Xz.shape[0])) # take up to 10 PCs or n_items
scores = pca.fit_transform(Xz) # (n_items, n_components)
expl_var = pca.explained_variance_ratio_

# Build tidy results frame
pc_cols = [f"PC{i+1}" for i in range(scores.shape[1])]
res_df = pd.DataFrame(scores, columns=pc_cols)
res_df.insert(0, "filename", filenames)
res_df["dif"] = dif_vec
res_df["dis"] = dis_vec

```

```

# ---- Correlate PCs with dif & dis ----
corr_rows = []
for j, pc in enumerate(pc_cols, start=1):
    r_dif, p_dif = pearsonr(res_df[pc], res_df["dif"])
    r_dis, p_dis = pearsonr(res_df[pc], res_df["dis"])
    corr_rows.append({
        "PC": f"PC{j}",
        "ExplainedVar(%)": round(100*expl_var[j-1], 2) if j-1 < len(expl_var) else np.nan,
        "r(PC, dif)": round(r_dif, 3), "p_dif": p_dif,
        "r(PC, dis)": round(r_dis, 3), "p_dis": p_dis
    })
corr_df = pd.DataFrame(corr_rows)

# ---- Print quick summary ----
print("== PCA on raw embeddings ==")
print(f"Items: {n_items}, Features: {X.shape[1]}")
print("\nTop PCs variance (%):", [round(100*v, 2) for v in expl_var[:5]])
print("\nCorrelations of PCs with dif and dis:")
display(corr_df)

# ---- Quick scatter diagnostics for top 3 PCs ----
k = min(3, scores.shape[1])
fig, axes = plt.subplots(k, 2, figsize=(10, 3.3*k))
axes = np.atleast_2d(axes)
for i in range(k):
    pc = pc_cols[i]
    r_dif, _ = pearsonr(res_df[pc], res_df["dif"])
    r_dis, _ = pearsonr(res_df[pc], res_df["dis"])
    # PC vs dif
    ax = axes[i, 0]
    ax.scatter(res_df[pc], res_df["dif"], c="tab:blue")
    ax.set_xlabel(pc); ax.set_ylabel("dif")
    ax.set_title(f"{pc} vs dif (r={r_dif:.2f})")
    ax.grid(True, alpha=.3)
    # PC vs dis
    ax = axes[i, 1]
    ax.scatter(res_df[pc], res_df["dis"], c="tab:orange")
    ax.set_xlabel(pc); ax.set_ylabel("dis")
    ax.set_title(f"{pc} vs dis (r={r_dis:.2f})")
    ax.grid(True, alpha=.3)

plt.tight_layout()
plt.show()

# ---- Optional: export the per-item PC scores & IRT params ----
# res_df.to_csv("pca_embedding_vs_irt.csv", index=False)

```

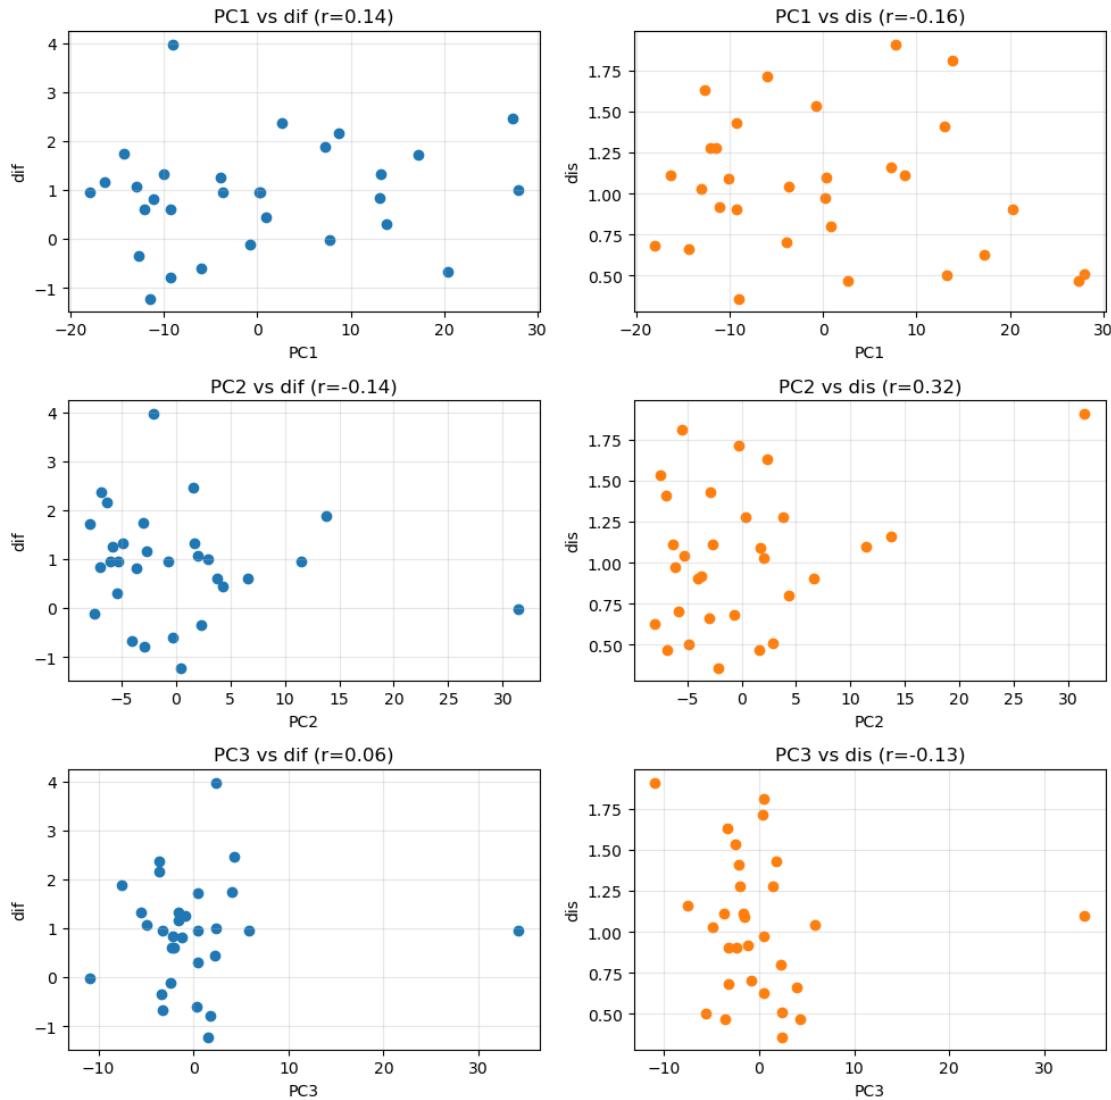
== PCA on raw embeddings ==

Items: 30, Features: 512

Top PCs variance (%): [32.05, 12.22, 10.3, 8.75, 5.56]

Correlations of PCs with dif and dis:

	PC	ExplainedVar(%)	r(PC, dif)	p_dif	r(PC, dis)	p_dis
0	PC1	32.05	0.141	0.458439	-0.164	0.387055
1	PC2	12.22	-0.139	0.462894	0.318	0.086680
2	PC3	10.30	0.057	0.765581	-0.131	0.489806
3	PC4	8.75	0.060	0.753443	-0.053	0.779854
4	PC5	5.56	0.010	0.959202	0.081	0.670957
5	PC6	4.50	0.065	0.733924	0.222	0.237748
6	PC7	3.18	0.198	0.295263	-0.107	0.573190
7	PC8	2.84	0.096	0.614234	-0.134	0.481732
8	PC9	2.33	-0.083	0.661331	-0.156	0.409502
9	PC10	2.10	-0.010	0.958012	-0.003	0.989374



```
[13]: """
Full Pseudo Factor Analysis for Matrix Reasoning Items
Uses multimodal LLM embeddings instead of CLIP
"""
```

```
import os
import glob
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from PIL import Image
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.decomposition import FactorAnalysis
```

```

from scipy.stats import pearsonr
import base64
from io import BytesIO

# =====
# CONFIGURATION - ENTER YOUR DETAILS HERE
# =====

# Choose your API: 'openai', 'anthropic', or 'google'
API_PROVIDER = 'openai'

# API Key - reads from environment variable
# Set in terminal before running:
#   Mac/Linux:   export OPENAI_API_KEY='your-key-here'
#   Windows:    set OPENAI_API_KEY=your-key-here
API_KEY = os.getenv('OPENAI_API_KEY') or os.getenv('ANTHROPIC_API_KEY') or os.getenv('GOOGLE_API_KEY')

# Image paths
IMAGE_FOLDER = './' # Folder containing your PNG files
IMAGE_PATTERN = '*.png' # Pattern to match images
MAX_IMAGES = 30 # Limit number of images

# IRT parameters CSV
IRT_CSV = 'form1.csv' # Should have columns: 'dis', 'dif'

# Number of factors for FA
N_FACTORS = 3

# =====
# HELPER FUNCTIONS
# =====

def encode_image_base64(image_path):
    """Convert image to base64 string"""
    with open(image_path, "rb") as f:
        return base64.b64encode(f.read()).decode('utf-8')

def get_embeddings_openai(image_paths, api_key):
    """Get embeddings from OpenAI GPT-4V"""
    from openai import OpenAI
    client = OpenAI(api_key=api_key)

    embeddings = []
    for img_path in image_paths:
        print(f"Processing {os.path.basename(img_path)}...")

```

```

# Encode image
base64_image = encode_image_base64(img_path)

# Get embedding via completion with special prompt
response = client.chat.completions.create(
    model="gpt-4o",
    messages=[{
        "role": "user",
        "content": [
            {"type": "text", "text": "Analyze this matrix reasoning ↵item. Describe its structure, complexity, and reasoning demands in detail."},
            {"type": "image_url", "image_url": {"url": f"data:image/png; ↵base64,{base64_image}"}}
        ]
    }],
    max_tokens=300
)

# Use text embedding of the response as proxy
text_response = response.choices[0].message.content
embed_response = client.embeddings.create(
    input=text_response,
    model="text-embedding-3-large"
)
embeddings.append(embed_response.data[0].embedding)

return np.array(embeddings)

def get_embeddings_anthropic(image_paths, api_key):
    """Get embeddings from Claude via description method"""
    import anthropic
    client = anthropic.Anthropic(api_key=api_key)

    embeddings = []
    for img_path in image_paths:
        print(f"Processing {os.path.basename(img_path)}...")

        base64_image = encode_image_base64(img_path)

        # Get detailed description
        message = client.messages.create(
            model="claude-3-5-sonnet-20241022",
            max_tokens=300,
            messages=[{
                "role": "user",
                "content": [
                    {

```

```

        "type": "image",
        "source": [
            "type": "base64",
            "media_type": "image/png",
            "data": base64_image,
        ],
    },
    {
        "type": "text",
        "text": "Analyze this matrix reasoning item. Describe:  

        ↵(1) number and types of objects, (2) spatial arrangement, (3) transformation  

        ↵rules, (4) cognitive complexity, (5) reasoning demands. Be specific and  

        ↵systematic."
    }
],
)
)

# Convert description to embedding using a sentence transformer
from sentence_transformers import SentenceTransformer
model = SentenceTransformer('all-mpnet-base-v2')
text_response = message.content[0].text
embedding = model.encode(text_response, convert_to_numpy=True)
embeddings.append(embedding)

return np.array(embeddings)

def get_embeddings_google(image_paths, api_key):
    """Get embeddings from Google Gemini"""
    import google.generativeai as genai
    genai.configure(api_key=api_key)

    model = genai.GenerativeModel('gemini-1.5-flash')

    embeddings = []
    for img_path in image_paths:
        print(f"Processing {os.path.basename(img_path)}...")

        img = Image.open(img_path)

        response = model.generate_content([
            "Analyze this matrix reasoning item systematically. Describe:  

            ↵structural elements, transformation patterns, cognitive demands, and  

            ↵problem-solving complexity.",
            img
        ])

```

```

# Convert description to embedding
from sentence_transformers import SentenceTransformer
embed_model = SentenceTransformer('all-mpnet-base-v2')
embedding = embed_model.encode(response.text, convert_to_numpy=True)
embeddings.append(embedding)

return np.array(embeddings)

# =====
# MAIN ANALYSIS
# =====

def run_full_analysis():
    print("=="*70)
    print("PSEUDO FACTOR ANALYSIS FOR MATRIX REASONING ITEMS")
    print("=="*70)

    # 1. Load images
    print("\n[1] Loading images...")
    image_paths = sorted(glob.glob(os.path.join(IMAGE_FOLDER, IMAGE_PATTERN)))[
        :MAX_IMAGES]
    print(f"Found {len(image_paths)} images")

    filenames = [os.path.basename(p) for p in image_paths]

    # 2. Load IRT parameters
    print("\n[2] Loading IRT parameters...")
    irt_df = pd.read_csv(IRT_CSV)
    dif = irt_df['dif'].values[:len(image_paths)]
    dis = irt_df['dis'].values[:len(image_paths)]
    print(f"Loaded {len(dif)} difficulty and {len(dis)} discrimination values")

    # 3. Get embeddings
    print(f"\n[3] Getting embeddings from {API_PROVIDER.upper()}...")

    if API_PROVIDER == 'openai':
        embeddings = get_embeddings_openai(image_paths, API_KEY)
    elif API_PROVIDER == 'anthropic':
        embeddings = get_embeddings_anthropic(image_paths, API_KEY)
    elif API_PROVIDER == 'google':
        embeddings = get_embeddings_google(image_paths, API_KEY)
    else:
        raise ValueError(f"Unknown API provider: {API_PROVIDER}")

    print(f"Embeddings shape: {embeddings.shape}")

    # 4. Compute cosine similarity matrix

```

```

print("\n[4] Computing similarity matrix...")
sim_matrix = cosine_similarity(embeddings)

# 5. Factor Analysis
print(f"\n[5] Running Factor Analysis (n_factors={N_FACTORS})...")
fa = FactorAnalysis(n_components=N_FACTORS, rotation=None)
fa.fit(sim_matrix)

# Get loadings
loadings = fa.components_.T
factor_scores = fa.transform(sim_matrix)

# 6. Create results dataframe
print("\n[6] Organizing results...")
results_df = pd.DataFrame(
    loadings,
    columns=[f'Factor{i+1}' for i in range(N_FACTORS)])
)
results_df.insert(0, 'filename', filenames)
results_df['dif'] = dif
results_df['dis'] = dis

# 7. Compute correlations
print("\n[7] Computing correlations between factors and IRT parameters...")
print("="*70)
print("CORRELATIONS: FACTORS vs IRT PARAMETERS")
print("="*70)

for i in range(N_FACTORS):
    factor_col = f'Factor{i+1}'
    r_dif, p_dif = pearsonr(results_df[factor_col], results_df['dif'])
    r_dis, p_dis = pearsonr(results_df[factor_col], results_df['dis'])

    print(f"\n{factor_col}:")
    print(f"  vs Difficulty:      r = {r_dif:+.3f}  (p = {p_dif:.4f})")
    print(f"  vs Discrimination: r = {r_dis:+.3f}  (p = {p_dis:.4f})")

# 8. Visualizations
print("\n[8] Creating visualizations...")

# Scree plot
component_vars = np.var(factor_scores, axis=0)
total_var = np.sum(component_vars)
explained_var_ratio = component_vars / total_var

fig, axes = plt.subplots(2, 2, figsize=(14, 10))

```

```

# Scree plot
axes[0, 0].plot(range(1, N_FACTORS + 1), explained_var_ratio, 'bo-', □
↳ linewidth=2)
axes[0, 0].set_xlabel('Factor', fontsize=11)
axes[0, 0].set_ylabel('Proportion of Variance', fontsize=11)
axes[0, 0].set_title('Scree Plot - Factor Analysis', fontsize=12, □
↳ fontweight='bold')
axes[0, 0].grid(True, alpha=0.3)

# Factor1 vs parameters
axes[0, 1].scatter(results_df['Factor1'], results_df['dif'], alpha=0.6, □
↳ s=60)
r_dif, _ = pearsonr(results_df['Factor1'], results_df['dif'])
axes[0, 1].set_xlabel('Factor 1', fontsize=11)
axes[0, 1].set_ylabel('Difficulty', fontsize=11)
axes[0, 1].set_title(f'Factor 1 vs Difficulty (r={r_dif:.3f})', fontsize=12)
axes[0, 1].grid(True, alpha=0.3)

axes[1, 0].scatter(results_df['Factor1'], results_df['dis'], alpha=0.6, □
↳ s=60, color='orange')
r_dis, _ = pearsonr(results_df['Factor1'], results_df['dis'])
axes[1, 0].set_xlabel('Factor 1', fontsize=11)
axes[1, 0].set_ylabel('Discrimination', fontsize=11)
axes[1, 0].set_title(f'Factor 1 vs Discrimination (r={r_dis:.3f})', □
↳ fontsize=12)
axes[1, 0].grid(True, alpha=0.3)

# Similarity matrix heatmap
im = axes[1, 1].imshow(sim_matrix, cmap='viridis', aspect='auto')
axes[1, 1].set_title('Cosine Similarity Matrix', fontsize=12, □
↳ fontweight='bold')
axes[1, 1].set_xlabel('Item', fontsize=11)
axes[1, 1].set_ylabel('Item', fontsize=11)
plt.colorbar(im, ax=axes[1, 1])

plt.tight_layout()
plt.savefig('pfa_results.png', dpi=150, bbox_inches='tight')
print("Saved: pfa_results.png")
plt.show()

# 9. Save results
print("\n[9] Saving results...")
results_df.to_csv('pfa_loadings_with_irt.csv', index=False)
print("Saved: pfa_loadings_with_irt.csv")

sim_df = pd.DataFrame(sim_matrix, columns=filenames, index=filenames)

```

```

sim_df.to_csv('similarity_matrix.csv')
print("Saved: similarity_matrix.csv")

print("\n" + "="*70)
print("ANALYSIS COMPLETE!")
print("="*70)

return results_df

# =====
# RUN
# =====

if __name__ == "__main__":
    # Check API key
    if API_KEY == 'your-api-key-here':
        print("ERROR: Please enter your API key in the API_KEY variable")
    else:
        results = run_full_analysis()
        print("\nResults preview:")
        print(results.head())

```

=====
PSEUDO FACTOR ANALYSIS FOR MATRIX REASONING ITEMS
=====

[1] Loading images...
Found 30 images

[2] Loading IRT parameters...
Loaded 30 difficulty and 30 discrimination values

[3] Getting embeddings from OPENAI...
Processing A-1.png...
Processing A-10.png...
Processing A-11.png...
Processing A-12.png...
Processing A-13.png...
Processing A-14.png...
Processing A-15.png...
Processing A-16.png...
Processing A-17.png...
Processing A-18.png...
Processing A-19.png...
Processing A-2.png...
Processing A-20.png...
Processing A-21.png...
Processing A-22.png...

```
Processing A-23.png...
Processing A-24.png...
Processing A-25.png...
Processing A-26.png...
Processing A-27.png...
Processing A-28.png...
Processing A-29.png...
Processing A-3.png...
Processing A-30.png...
Processing A-4.png...
Processing A-5.png...
Processing A-6.png...
Processing A-7.png...
Processing A-8.png...
Processing A-9.png...
Embeddings shape: (30, 3072)
```

[4] Computing similarity matrix...

[5] Running Factor Analysis (n_factors=3)...

[6] Organizing results...

[7] Computing correlations between factors and IRT parameters...

```
=====
CORRELATIONS: FACTORS vs IRT PARAMETERS
=====
```

Factor1:

```
    vs Difficulty:      r = -0.042  (p = 0.8255)
    vs Discrimination: r = +0.055  (p = 0.7740)
```

Factor2:

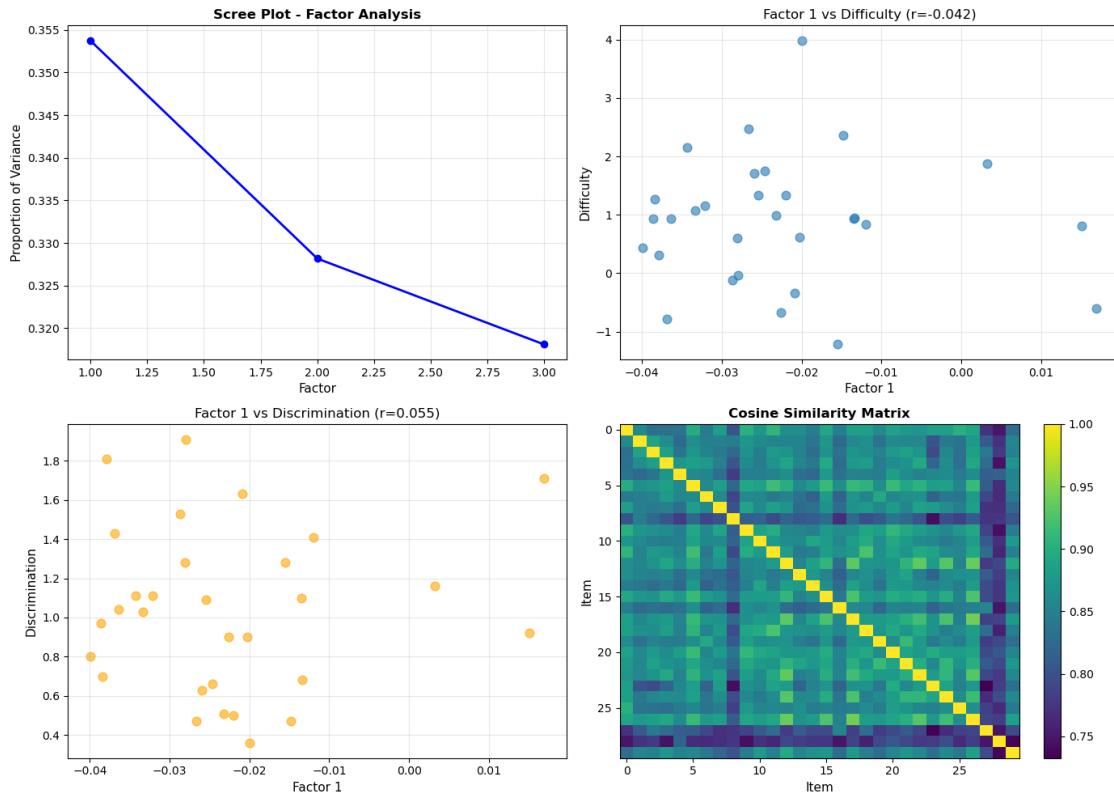
```
    vs Difficulty:      r = -0.102  (p = 0.5933)
    vs Discrimination: r = -0.029  (p = 0.8810)
```

Factor3:

```
    vs Difficulty:      r = -0.036  (p = 0.8522)
    vs Discrimination: r = +0.073  (p = 0.7025)
```

[8] Creating visualizations...

Saved: pfa_results.png



```
[9] Saving results...
Saved: pfa_loadings_with_irt.csv
Saved: similarity_matrix.csv
```

```
=====
ANALYSIS COMPLETE!
=====
```

Results preview:

	filename	Factor1	Factor2	Factor3	dif	dis
0	A-1.png	-0.032143	-0.019830	0.009578	1.15	1.11
1	A-10.png	-0.013488	-0.005540	-0.018043	0.94	1.10
2	A-11.png	-0.015542	-0.000804	-0.023565	-1.22	1.28
3	A-12.png	-0.025428	0.001947	-0.021314	1.33	1.09
4	A-13.png	-0.013348	0.001396	0.003446	0.95	0.68

```
[21]: # let's try predicting from CLIP for dif and dis
```

```
"""

```

*Comprehensive prediction of IRT parameters from CLIP embeddings
Includes: Ridge, Lasso, ElasticNet, Random Forest, Gradient Boosting, XGBoost*

All with proper cross-validation and feature resampling

```

"""
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import cross_val_score, KFold, GridSearchCV
from sklearn.linear_model import Ridge, Lasso, ElasticNet
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from scipy.stats import pearsonr
import warnings
warnings.filterwarnings('ignore')

# Optional: XGBoost (install with: pip install xgboost)
try:
    from xgboost import XGBRegressor
    HAS_XGBOOST = True
except ImportError:
    HAS_XGBOOST = False
    print("Note: XGBoost not installed. Install with: pip install xgboost")

# =====
# ASSUMES YOU ALREADY HAVE THESE FROM YOUR NOTEBOOK:
# - embeddings: np.array shape (30, 512) - CLIP embeddings
# - dif: np.array shape (30,) - difficulty parameters
# - dis: np.array shape (30,) - discrimination parameters
# =====

# If running standalone, uncomment and load:
# embeddings = np.load('clip_embeddings.npy')
# irt_df = pd.read_csv('form1.csv')
# dif = irt_df['dif'].values
# dis = irt_df['dis'].values

def run_comprehensive_regression(X, y_dif, y_dis, target_name='dif'):
    """
    Run multiple regression models with cross-validation

    Parameters:
    -----
    X : array-like, shape (n_items, n_features)
        CLIP embeddings (typically 30 x 512)
    y_dif : array-like, shape (n_items,)
        Difficulty parameters
    y_dis : array-like, shape (n_items,)
        Discrimination parameters
    """

    # Import necessary modules
    import numpy as np
    import pandas as pd
    import matplotlib.pyplot as plt
    from sklearn.model_selection import cross_val_score, KFold, GridSearchCV
    from sklearn.linear_model import Ridge, Lasso, ElasticNet
    from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
    from sklearn.preprocessing import StandardScaler
    from sklearn.pipeline import Pipeline
    from scipy.stats import pearsonr
    import warnings
    warnings.filterwarnings('ignore')

    # Check if XGBoost is available
    try:
        from xgboost import XGBRegressor
        HAS_XGBOOST = True
    except ImportError:
        HAS_XGBOOST = False
        print("Note: XGBoost not installed. Install with: pip install xgboost")

    # Load data if running standalone
    # embeddings = np.load('clip_embeddings.npy')
    # irt_df = pd.read_csv('form1.csv')
    # dif = irt_df['dif'].values
    # dis = irt_df['dis'].values

    # Create pipeline
    pipeline = Pipeline([
        ('scaler', StandardScaler()),
        ('regressor', Ridge())
    ])

    # Define cross-validation
    kf = KFold(n_splits=5)

    # Define scoring metric
    def neg_mean_squared_error(y_true, y_pred):
        return -np.mean((y_true - y_pred) ** 2)

    # Perform cross-validation
    scores = cross_val_score(pipeline, X, y_dif, cv=kf, scoring=neg_mean_squared_error)

    # Print results
    print(f"Ridge Regression scores: {scores}")
    print(f"Mean score: {np.mean(scores)}")
    print(f"Standard deviation: {np.std(scores)}")

    # Check if XGBoost is available
    if HAS_XGBOOST:
        # Create XGBoost pipeline
        xgb_pipeline = Pipeline([
            ('scaler', StandardScaler()),
            ('xgb', XGBRegressor())
        ])

        # Perform cross-validation
        xgb_scores = cross_val_score(xgb_pipeline, X, y_dif, cv=kf, scoring=neg_mean_squared_error)

        # Print results
        print(f"\nXGBoost scores: {xgb_scores}")
        print(f"Mean score: {np.mean(xgb_scores)}")
        print(f"Standard deviation: {np.std(xgb_scores)}")

    # Plot results
    plt.figure(figsize=(10, 6))
    plt.title("Comparison of Ridge Regression and XGBoost")
    plt.bar(['Ridge', 'XGBoost'], [np.mean(scores), np.mean(xgb_scores)])
    plt.xlabel("Model")
    plt.ylabel("Mean Squared Error")
    plt.show()

```

```

Discrimination parameters
target_name : str
    Which target to predict ('dif' or 'dis')
"""

y = y_dif if target_name == 'dif' else y_dis
n_items, n_features = X.shape

print("=="*80)
print(f"PREDICTING {target_name.upper()} FROM CLIP EMBEDDINGS")
print("=="*80)
print(f"Items: {n_items}, Features: {n_features}")
print(f"Target range: [{y.min():.2f}, {y.max():.2f}]")
print()

# Cross-validation setup
# With 30 items, use 5-fold CV (6 items per fold)
cv = KFold(n_splits=5, shuffle=True, random_state=42)

# Store results
results = []

# =====
# 1. RIDGE REGRESSION (L2 regularization)
# =====

print("[1/6] Ridge Regression (L2 penalty)...")

ridge_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('ridge', Ridge())
])

ridge_params = {
    'ridge__alpha': [0.01, 0.1, 1.0, 10.0, 100.0, 1000.0]
}

ridge_grid = GridSearchCV(
    ridge_pipeline, ridge_params,
    cv=cv, scoring='r2', n_jobs=-1
)
ridge_grid.fit(X, y)

ridge_cv_scores = cross_val_score(
    ridge_grid.best_estimator_, X, y, cv=cv, scoring='r2'
)

results.append({

```

```

        'Model': 'Ridge',
        'Best_Parms': f"={ridge_grid.best_params_['ridge_alpha']}",
        'CV_R2_mean': ridge_cv_scores.mean(),
        'CV_R2_std': ridge_cv_scores.std(),
        'Best_Model': ridge_grid.best_estimator_
    })

    print(f"  Best : {ridge_grid.best_params_['ridge_alpha']}")
    print(f"  CV R2: {ridge_cv_scores.mean():.3f} ± {ridge_cv_scores.std():.3f}")
print()

# =====
# 2. LASSO REGRESSION (L1 regularization, feature selection)
# =====
print("[2/6] Lasso Regression (L1 penalty, sparse)...")

lasso_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('lasso', Lasso(max_iter=5000))
])

lasso_params = {
    'lasso_alpha': [0.001, 0.01, 0.1, 1.0, 10.0]
}

lasso_grid = GridSearchCV(
    lasso_pipeline, lasso_params,
    cv=cv, scoring='r2', n_jobs=-1
)
lasso_grid.fit(X, y)

lasso_cv_scores = cross_val_score(
    lasso_grid.best_estimator_, X, y, cv=cv, scoring='r2'
)

# Count non-zero coefficients (feature selection)
lasso_grid.best_estimator_.fit(X, y)
n_selected = np.sum(lasso_grid.best_estimator_.named_steps['lasso'].coef_ != 0)

results.append({
    'Model': 'Lasso',
    'Best_Parms': f"={lasso_grid.best_params_['lasso_alpha']},\n{n_selected} features",
    'CV_R2_mean': lasso_cv_scores.mean(),
    'CV_R2_std': lasso_cv_scores.std(),
})

```

```

        'Best_Model': lasso_grid.best_estimator_
    })

    print(f"  Best : {lasso_grid.best_params_['lasso_alpha']}") 
    print(f"  Selected features: {n_selected}/{n_features}")
    print(f"  CV R2: {lasso_cv_scores.mean():.3f} ± {lasso_cv_scores.std():.
        ↪3f}")
    print()

# =====
# 3. ELASTIC NET (L1 + L2 regularization)
# =====

print("[3/6] ElasticNet (L1 + L2 penalty)...")

enet_pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('enet', ElasticNet(max_iter=5000))
])

enet_params = {
    'enet__alpha': [0.01, 0.1, 1.0, 10.0],
    'enet__l1_ratio': [0.1, 0.5, 0.9]  # 0=Ridge, 1=Lasso
}

enet_grid = GridSearchCV(
    enet_pipeline, enet_params,
    cv=cv, scoring='r2', n_jobs=-1
)
enet_grid.fit(X, y)

enet_cv_scores = cross_val_score(
    enet_grid.best_estimator_, X, y, cv=cv, scoring='r2'
)

results.append({
    'Model': 'ElasticNet',
    'Best_Params': f"={enet_grid.best_params_['enet_alpha']},\n        ↪l1={enet_grid.best_params_['enet_l1_ratio']}",
    'CV_R2_mean': enet_cv_scores.mean(),
    'CV_R2_std': enet_cv_scores.std(),
    'Best_Model': enet_grid.best_estimator_
})

print(f"  Best : {enet_grid.best_params_['enet_alpha']}")
print(f"  Best l1_ratio: {enet_grid.best_params_['enet_l1_ratio']}")
print(f"  CV R2: {enet_cv_scores.mean():.3f} ± {enet_cv_scores.std():.3f}")
print()

```

```

# =====
# 4. RANDOM FOREST (Bootstrap aggregating with feature resampling)
# =====
print("[4/6] Random Forest (Bootstrap + Feature Resampling)...")

# Random Forest naturally does bootstrap resampling (bootstrap=True by default)
# and feature resampling (max_features < n_features)

rf_params = {
    'n_estimators': [100, 200, 500],
    'max_depth': [3, 5, 7, None],
    'min_samples_split': [2, 5, 10],
    'max_features': ['sqrt', 'log2', 0.3] # Feature resampling
}

rf = RandomForestRegressor(random_state=42, n_jobs=-1)
rf_grid = GridSearchCV(
    rf, rf_params,
    cv=cv, scoring='r2', n_jobs=-1
)
rf_grid.fit(X, y)

rf_cv_scores = cross_val_score(
    rf_grid.best_estimator_, X, y, cv=cv, scoring='r2'
)

results.append({
    'Model': 'RandomForest',
    'Best_Params': f"trees={rf_grid.best_params_['n_estimators']}, depth={rf_grid.best_params_['max_depth']}",
    'CV_R2_mean': rf_cv_scores.mean(),
    'CV_R2_std': rf_cv_scores.std(),
    'Best_Model': rf_grid.best_estimator_
})

print(f" Best n_estimators: {rf_grid.best_params_['n_estimators']} ")
print(f" Best max_depth: {rf_grid.best_params_['max_depth']} ")
print(f" Best max_features: {rf_grid.best_params_['max_features']} ")
print(f" CV R2: {rf_cv_scores.mean():.3f} ± {rf_cv_scores.std():.3f} ")
print()

# =====
# 5. GRADIENT BOOSTING (Sequential boosting)
# =====
print("[5/6] Gradient Boosting (Sequential Boosting)...")

```

```

gb_params = {
    'n_estimators': [100, 200, 500],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0] # Stochastic boosting (resampling)
}

gb = GradientBoostingRegressor(random_state=42)
gb_grid = GridSearchCV(
    gb, gb_params,
    cv=cv, scoring='r2', n_jobs=-1
)
gb_grid.fit(X, y)

gb_cv_scores = cross_val_score(
    gb_grid.best_estimator_, X, y, cv=cv, scoring='r2'
)

results.append({
    'Model': 'GradientBoosting',
    'Best_Params': f"trees={gb_grid.best_params_['n_estimators']}",
    ↴lr={gb_grid.best_params_['learning_rate']},
    'CV_R²_mean': gb_cv_scores.mean(),
    'CV_R²_std': gb_cv_scores.std(),
    'Best_Model': gb_grid.best_estimator_
})

print(f" Best n_estimators: {gb_grid.best_params_['n_estimators']} ")
print(f" Best learning_rate: {gb_grid.best_params_['learning_rate']} ")
print(f" Best subsample: {gb_grid.best_params_['subsample']} ")
print(f" CV R²: {gb_cv_scores.mean():.3f} ± {gb_cv_scores.std():.3f} ")
print()

# =====#
# 6. XGBOOST (Advanced boosting with regularization)
# =====#
if HAS_XGBOOST:
    print("[6/6] XGBoost (Regularized Boosting)...")

xgb_params = {
    'n_estimators': [100, 200, 500],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.1, 0.2],
    'subsample': [0.8, 1.0],
    'colsample_bytree': [0.8, 1.0], # Feature resampling per tree
    'reg_alpha': [0, 0.1, 1.0], # L1 regularization
}

```

```

        'reg_lambda': [1.0, 10.0] # L2 regularization
    }

xgb = XGBRegressor(random_state=42, n_jobs=-1)
xgb_grid = GridSearchCV(
    xgb, xgb_params,
    cv=cv, scoring='r2', n_jobs=-1
)
xgb_grid.fit(X, y)

xgb_cv_scores = cross_val_score(
    xgb_grid.best_estimator_, X, y, cv=cv, scoring='r2'
)

results.append({
    'Model': 'XGBoost',
    'Best_Params': f"trees={xgb_grid.best_params_['n_estimators']},",
    ↪lr={xgb_grid.best_params_['learning_rate']}",
    'CV_R2_mean': xgb_cv_scores.mean(),
    'CV_R2_std': xgb_cv_scores.std(),
    'Best_Model': xgb_grid.best_estimator_
})

print(f" Best n_estimators: {xgb_grid.best_params_['n_estimators']} ")
print(f" Best learning_rate: {xgb_grid.best_params_['learning_rate']} ")
print(f" CV R2: {xgb_cv_scores.mean():.3f} ± {xgb_cv_scores.std():.
↪3f}")
print()
else:
    print("[6/6] XGBoost skipped (not installed)")
    print()

# =====#
# SUMMARY
# =====#
results_df = pd.DataFrame(results)
results_df = results_df.sort_values('CV_R2_mean', ascending=False)

print("=*80)
print(f"RESULTS SUMMARY - Predicting {target_name.upper()}")
print("=*80)
print(results_df[['Model', 'CV_R2_mean', 'CV_R2_std', 'Best_Params']].
↪to_string(index=False))
print()

best_model_name = results_df.iloc[0]['Model']
best_r2 = results_df.iloc[0]['CV_R2_mean']

```

```

best_model = results_df.iloc[0]['Best_Model']

print(f" BEST MODEL: {best_model_name} (CV R2 = {best_r2:.3f})")
print()

# =====
# VISUALIZATION
# =====

fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Bar plot of CV R2
ax = axes[0]
models = results_df['Model'].values
r2_means = results_df['CV_R2_mean'].values
r2_stds = results_df['CV_R2_std'].values

colors = ['#1f77b4' if m != best_model_name else '#ff7f0e' for m in models]
bars = ax.barh(models, r2_means, xerr=r2_stds, color=colors, alpha=0.8)
ax.set_xlabel('Cross-Validated R2', fontsize=12)
ax.set_title(f'Model Comparison - Predicting {target_name.upper()}', fontweight='bold')
ax.axvline(0, color='black', linewidth=0.8)
ax.grid(axis='x', alpha=0.3)

# Predicted vs Actual for best model
ax = axes[1]
best_model.fit(X, y)
y_pred = best_model.predict(X)

ax.scatter(y, y_pred, alpha=0.6, s=80, edgecolors='black', linewidth=0.5)

# Perfect prediction line
min_val, max_val = y.min(), y.max()
ax.plot([min_val, max_val], [min_val, max_val], 'r--', linewidth=2, label='Perfect prediction')

# Correlation
r, p = pearsonr(y, y_pred)
ax.text(0.05, 0.95, f'r = {r:.3f}\nR2 = {best_r2:.3f}', transform=ax.transAxes, fontsize=11,
        verticalalignment='top', bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

ax.set_xlabel(f'Actual {target_name.upper()}', fontsize=12)
ax.set_ylabel(f'Predicted {target_name.upper()}', fontsize=12)
ax.set_title(f'{best_model_name} - Predicted vs Actual', fontweight='bold', fontsize=13)

```

```

    ax.legend()
    ax.grid(True, alpha=0.3)

    plt.tight_layout()
    plt.savefig(f'regression_results_{target_name}.png', dpi=150, u
    ↪bbox_inches='tight')
    print(f"Saved: regression_results_{target_name}.png")
    plt.show()

    return results_df, best_model

# =====
# RUN ANALYSIS
# =====

if __name__ == "__main__":
    # Load/check required data
    print("*"*80)
    print("LOADING DATA")
    print("*"*80)

    # Check if embeddings exist
    try:
        print(f" Embeddings found: {embeddings.shape}")
    except NameError:
        print(" ERROR: 'embeddings' not defined")
        print(" Re-run your CLIP encoding cell first!")
        import sys
        sys.exit(1)

    # Load IRT parameters from CSV
    try:
        import pandas as pd
        form1 = pd.read_csv('form1.csv')
        dif = form1['dif'].values
        dis = form1['dis'].values
        print(f" Loaded IRT parameters: {len(dif)} items")
        print(f" Difficulty range: [{dif.min():.2f}, {dif.max():.2f}]")
        print(f" Discrimination range: [{dis.min():.2f}, {dis.max():.2f}]")
    except FileNotFoundError:
        print(" ERROR: 'form1.csv' not found")
        print(" Make sure form1.csv is in the current directory")
        import sys
        sys.exit(1)
    except Exception as e:
        print(f" ERROR loading IRT parameters: {e}")
        import sys

```

```

    sys.exit(1)

print()

# Predict difficulty
print("\n" + "="*80)
print("PART 1: PREDICTING DIFFICULTY")
print("="*80 + "\n")
results_dif, best_model_dif = run_comprehensive_regression(
    embeddings, dif, dis, target_name='dif'
)

# Predict discrimination
print("\n" + "="*80)
print("PART 2: PREDICTING DISCRIMINATION")
print("="*80 + "\n")
results_dis, best_model_dis = run_comprehensive_regression(
    embeddings, dif, dis, target_name='dis'
)

# Final summary
print("\n" + "="*80)
print("FINAL SUMMARY")
print("="*80)
print(f"\nBest for DIFFICULTY: {results_dif.iloc[0]['Model']} "
      f"(CV R2 = {results_dif.iloc[0]['CV_R2_mean']:.3f})")
print(f"Best for DISCRIMINATION: {results_dis.iloc[0]['Model']} "
      f"(CV R2 = {results_dis.iloc[0]['CV_R2_mean']:.3f})")
print()

```

Note: XGBoost not installed. Install with: pip install xgboost

LOADING DATA

```

Embeddings found: (30, 512)
Loaded IRT parameters: 30 items
Difficulty range: [-1.22, 3.98]
Discrimination range: [0.36, 1.91]

```

PART 1: PREDICTING DIFFICULTY

PREDICTING DIF FROM CLIP EMBEDDINGS

Items: 30, Features: 512

Target range: [-1.22, 3.98]

[1/6] Ridge Regression (L2 penalty)...

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

To disable this warning, you can either:

- Avoid using `tokenizers` before the fork if possible
- Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)

huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...

```

To disable this warning, you can either:
  - Avoid using `tokenizers` before the fork if possible
  - Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)
huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...
To disable this warning, you can either:
  - Avoid using `tokenizers` before the fork if possible
  - Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)
huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...
To disable this warning, you can either:
  - Avoid using `tokenizers` before the fork if possible
  - Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)
huggingface/tokenizers: The current process just got forked, after parallelism has already been used. Disabling parallelism to avoid deadlocks...
To disable this warning, you can either:
  - Avoid using `tokenizers` before the fork if possible
  - Explicitly set the environment variable TOKENIZERS_PARALLELISM=(true | false)
Best : 1000.0
CV R2: -0.884 ± 0.721

```

[2/6] Lasso Regression (L1 penalty, sparse)...

```

Best : 1.0
Selected features: 0/512
CV R2: -0.562 ± 0.514

```

[3/6] ElasticNet (L1 + L2 penalty)...

```

/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 4.330e-03, tolerance: 3.115e-03
    model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.534e-03, tolerance: 1.955e-03
    model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 2.534e-03, tolerance: 1.955e-03

```

```

gap: 3.294e-03, tolerance: 3.115e-03
    model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.011e-02, tolerance: 3.115e-03
    model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 4.754e-03, tolerance: 2.704e-03
    model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 3.401e-03, tolerance: 1.955e-03
    model = cd_fast.enet_coordinate_descent(
/opt/anaconda3/lib/python3.11/site-
packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 1.565e-02, tolerance: 3.016e-03
    model = cd_fast.enet_coordinate_descent(
        Best : 1.0
        Best l1_ratio: 0.9
        CV R2: -0.562 ± 0.514

[4/6] Random Forest (Bootstrap + Feature Resampling)...
        Best n_estimators: 500
        Best max_depth: 5
        Best max_features: log2
        CV R2: -0.806 ± 0.706

[5/6] Gradient Boosting (Sequential Boosting)...
        Best n_estimators: 100
        Best learning_rate: 0.01
        Best subsample: 0.8
        CV R2: -0.892 ± 0.878

[6/6] XGBoost skipped (not installed)
=====
```

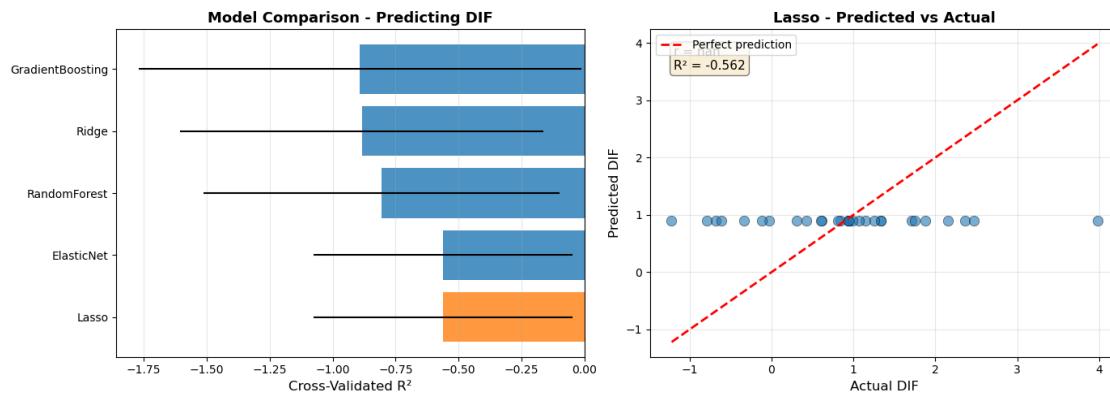
RESULTS SUMMARY – Predicting DIF

=====

Model	CV_R ² _mean	CV_R ² _std	Best_Params
Lasso	-0.561683	0.514341	=1.0, 0 features
ElasticNet	-0.561683	0.514341	=1.0, l1=0.9
RandomForest	-0.806240	0.705928	trees=500, depth=5
Ridge	-0.884478	0.720942	=1000.0
GradientBoosting	-0.892150	0.878359	trees=100, lr=0.01

BEST MODEL: Lasso (CV R² = -0.562)

Saved: regression_results_dif.png



PART 2: PREDICTING DISCRIMINATION

PREDICTING DIS FROM CLIP EMBEDDINGS

Items: 30, Features: 512

Target range: [0.36, 1.91]

[1/6] Ridge Regression (L2 penalty)...

Best : 1000.0

CV R²: -1.466 ± 1.687

[2/6] Lasso Regression (L1 penalty, sparse)...

Best : 1.0

Selected features: 0/512

CV R²: -1.047 ± 1.431

[3/6] ElasticNet (L1 + L2 penalty)...

/opt/anaconda3/lib/python3.11/site-

```

packages/sklearn/linear_model/_coordinate_descent.py:631: ConvergenceWarning:
Objective did not converge. You might want to increase the number of iterations,
check the scale of the features or consider increasing regularisation. Duality
gap: 5.250e-04, tolerance: 3.660e-04
    model = cd_fast.enet_coordinate_descent(
        Best : 1.0
        Best l1_ratio: 0.5
        CV R2: -1.047 ± 1.431

[4/6] Random Forest (Bootstrap + Feature Resampling)...
    Best n_estimators: 100
    Best max_depth: 3
    Best max_features: log2
    CV R2: -1.414 ± 1.640

[5/6] Gradient Boosting (Sequential Boosting)...
    Best n_estimators: 100
    Best learning_rate: 0.01
    Best subsample: 0.8
    CV R2: -1.682 ± 1.971

[6/6] XGBoost skipped (not installed)

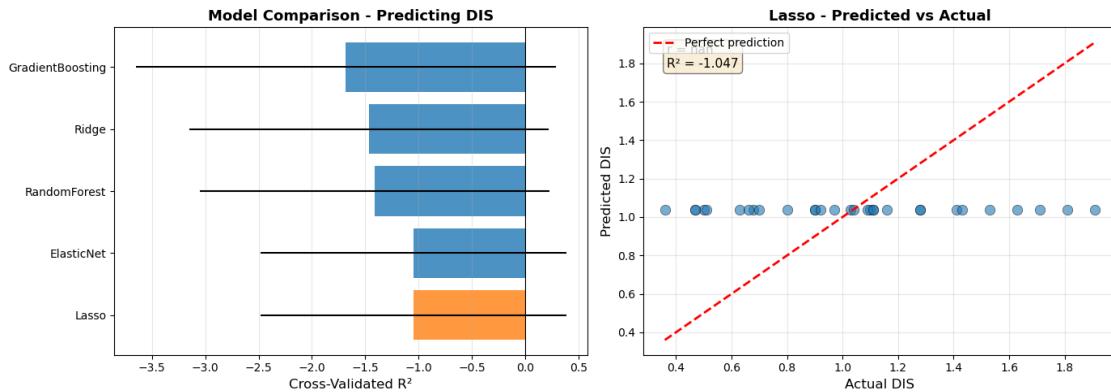
```

RESULTS SUMMARY – Predicting DIS

Model	CV_R ² _mean	CV_R ² _std	Best_Params
Lasso	-1.046656	1.431416	=1.0, 0 features
ElasticNet	-1.046656	1.431416	=1.0, l1=0.5
RandomForest	-1.413974	1.640421	trees=100, depth=3
Ridge	-1.466376	1.687025	=1000.0
GradientBoosting	-1.682395	1.970836	trees=100, lr=0.01

BEST MODEL: Lasso (CV R² = -1.047)

Saved: regression_results_dis.png



FINAL SUMMARY

Best for DIFFICULTY: Lasso (CV $R^2 = -0.562$)
 Best for DISCRIMINATION: Lasso (CV $R^2 = -1.047$)

```
[27]: # ===== TRY 1: DINOV2 Embeddings + Factor Analysis (FIXED) =====
import torch
import numpy as np
import pandas as pd
from PIL import Image
from transformers import AutoImageProcessor, AutoModel
from factor_analyzer import FactorAnalyzer
from factor_analyzer.factor_analyzer import calculate_bartlett_sphericity, calculate_kmo
import matplotlib.pyplot as plt
from scipy.stats import pearsonr

print("Loading DINOV2 model...")
processor = AutoImageProcessor.from_pretrained('facebook/dinov2-base')
model = AutoModel.from_pretrained('facebook/dinov2-base')
model.eval()

# Extract embeddings
print("Extracting DINOV2 embeddings...")
dino_embeddings = []
for img_path in image_paths: # Assumes you have image_paths list
    img = Image.open(img_path).convert('RGB')
    inputs = processor(images=img, return_tensors="pt")
    with torch.no_grad():


```

```

        outputs = model(**inputs)
        # Use [CLS] token (first token)
        embedding = outputs.last_hidden_state[:, 0, :].squeeze().numpy()
        dino_embeddings.append(embedding)

dino_embeddings = np.array(dino_embeddings)
print(f"DINOv2 embeddings shape: {dino_embeddings.shape}")

# Standardize
from sklearn.preprocessing import StandardScaler
dino_scaled = StandardScaler().fit_transform(dino_embeddings)

# Check factorability
print("\n--- Factorability Tests ---")
chi_square, p_value = calculate_bartlett_sphericity(dino_scaled)
print(f"Bartlett's test: {chi_square:.2f}, p={p_value:.4f}")
kmo_all, kmo_model = calculate_kmo(dino_scaled)
print(f"KMO: {kmo_model:.3f}")

# Factor Analysis with 3 factors
print("\n--- Factor Analysis (3 factors, ML) ---")
fa = FactorAnalyzer(n_factors=3, rotation='varimax', method='ml')
fa.fit(dino_scaled)

# Get loadings - THIS IS THE FIX
loadings = fa.loadings_ # This is (768, 3) - features x factors
# We need (30, 3) - items x factors, so we get factor scores instead
factor_scores = fa.transform(dino_scaled) # This gives (30, 3)

loadings_df = pd.DataFrame(
    factor_scores, # CHANGED: Use factor scores, not loadings
    columns=[f'Factor{i+1}' for i in range(3)],
    index=[f"item_{i+1}" for i in range(len(dino_embeddings))]
)

# Add IRT parameters
loadings_df['dif'] = form1['dif'].values # CHANGED: Use form1
loadings_df['dis'] = form1['dis'].values

print("\nFactor scores (first 10 items):")
print(loadings_df.head(10))

# Correlate factors with IRT parameters
print("\n--- Correlations with IRT parameters ---")
for factor in ['Factor1', 'Factor2', 'Factor3']:
    r_dif, p_dif = pearsonr(loadings_df[factor], loadings_df['dif'])
    r_dis, p_dis = pearsonr(loadings_df[factor], loadings_df['dis'])

```

```

print(f"factor:{factor}")
print(f" r(factor, dif) = {r_dif:.3f}, p = {p_dif:.4f}")
print(f" r(factor, dis) = {r_dis:.3f}, p = {p_dis:.4f}")

# Visualize
fig, axes = plt.subplots(3, 2, figsize=(10, 12))
for i, factor in enumerate(['Factor1', 'Factor2', 'Factor3']):
    # Factor vs difficulty
    axes[i, 0].scatter(loadings_df[factor], loadings_df['dif'], alpha=0.6)
    axes[i, 0].set_xlabel(factor)
    axes[i, 0].set_ylabel('Difficulty')
    r_dif, _ = pearsonr(loadings_df[factor], loadings_df['dif'])
    axes[i, 0].set_title(f'{factor} vs Difficulty (r={r_dif:.2f})')
    axes[i, 0].grid(True, alpha=0.3)

    # Factor vs discrimination
    axes[i, 1].scatter(loadings_df[factor], loadings_df['dis'], alpha=0.6)
    axes[i, 1].set_xlabel(factor)
    axes[i, 1].set_ylabel('Discrimination')
    r_dis, _ = pearsonr(loadings_df[factor], loadings_df['dis'])
    axes[i, 1].set_title(f'{factor} vs Discrimination (r={r_dis:.2f})')
    axes[i, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

print("\n DIN0v2 + FA complete!")

```

Loading DIN0v2 model...
Extracting DIN0v2 embeddings...
DIN0v2 embeddings shape: (30, 768)

--- Factorability Tests ---
Bartlett's test: $\chi^2=\text{inf}$, $p=1.0000$
KMO: nan

--- Factor Analysis (3 factors, ML) ---

Factor scores (first 10 items):

	Factor1	Factor2	Factor3	dif	dis
item_1	-0.288610	1.114952	-1.603965	1.15	1.11
item_2	-0.918159	1.782717	-1.276086	0.94	1.10
item_3	-0.731592	-1.623617	0.274825	-1.22	1.28
item_4	1.132702	-1.026272	-1.641842	1.33	1.09
item_5	1.613172	-1.482189	-1.665563	0.95	0.68
item_6	-1.982107	-0.617042	-0.632650	1.07	1.03
item_7	0.486858	-0.408625	0.075851	0.61	0.90
item_8	0.015957	0.838686	0.071655	0.60	1.28

```
item_9    0.212099  0.754642 -1.844838  1.88  1.16  
item_10   0.436989  0.363259  1.370189 -0.03  1.91
```

--- Correlations with IRT parameters ---

Factor1:

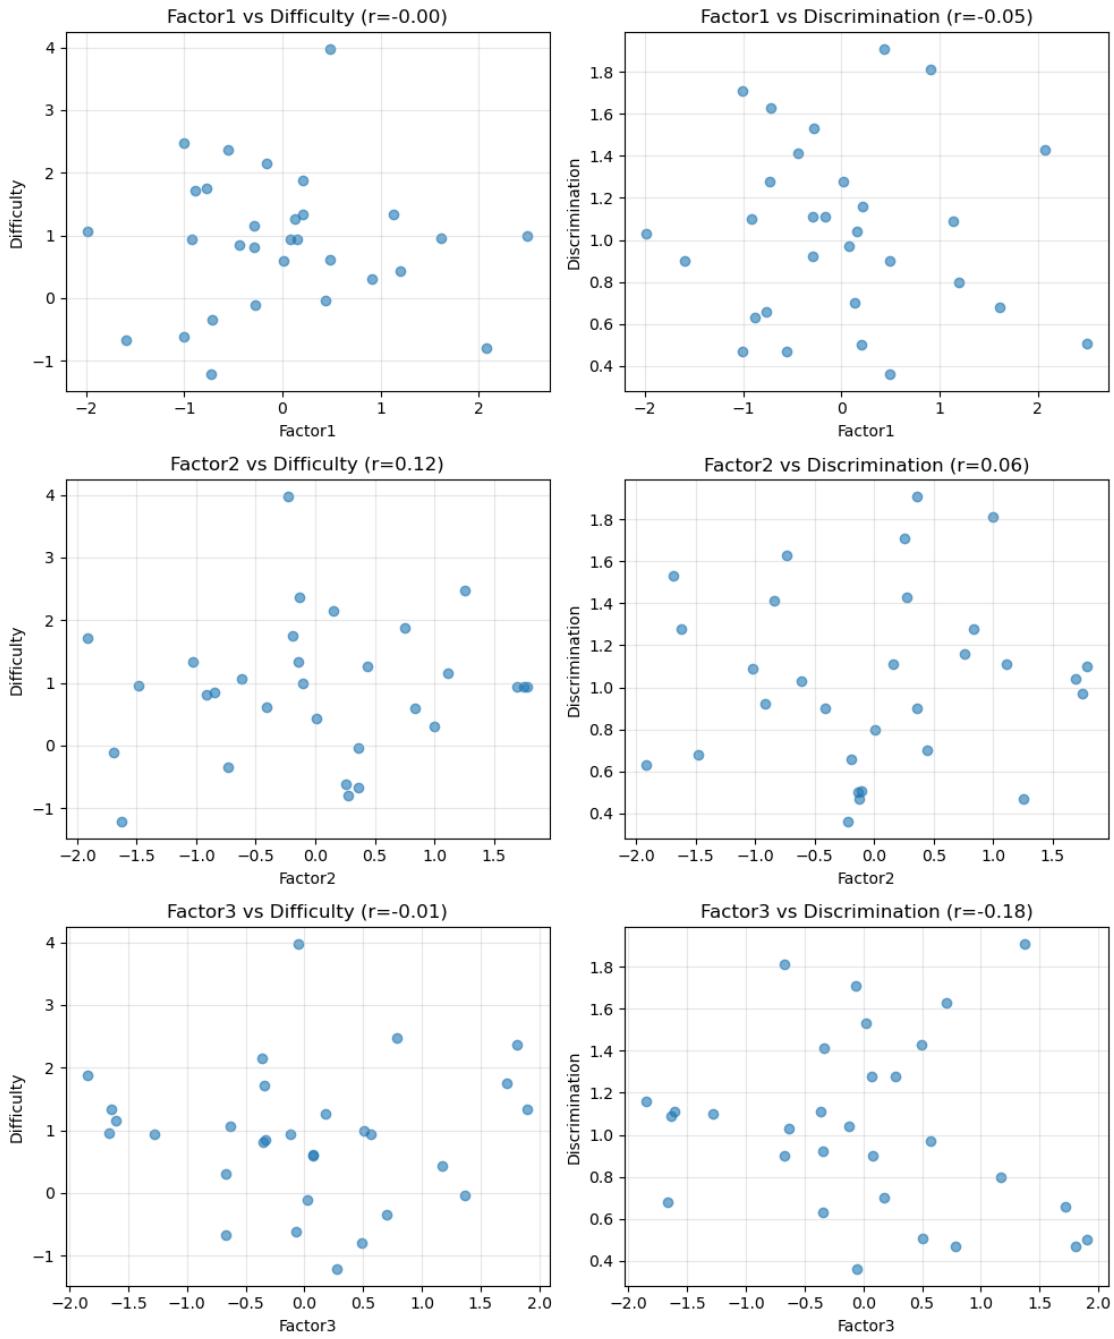
```
r(factor, dif) = -0.003, p = 0.9857  
r(factor, dis) = -0.049, p = 0.7968
```

Factor2:

```
r(factor, dif) = 0.123, p = 0.5169  
r(factor, dis) = 0.058, p = 0.7627
```

Factor3:

```
r(factor, dif) = -0.011, p = 0.9557  
r(factor, dis) = -0.177, p = 0.3488
```



DINOv2 + FA complete!

[]: